

ML Engineer

Project 4

E-commerce customers clustering



Victor Benard

MISSION STATEMENT

olist



Olist is a Brazilian online e-commerce platform that wishes to carry out a segmentation of its customers. Objective is to create clusters, and subsequently propose a maintenance plan based on clustering temporal stability.

TABLE OF CONTENTS

01 Overview and
targets

02 Data merging
and analysis

03 Clustering

04 Conclusion

01

OVERVIEW & TARGETS



OVERVIEW & TARGETS

Olist proposes an anonymized database that contains orders information and history, ordered products, satisfaction, and location of clients from 2017.

Objective:

Use unsupervised methods to regroup customers that have similar profiles.

Key information:

- Data available from October 2016 to October 2018
- 9 Dataframes for a total of 99k customers id
 - Customer, Seller, Geolocation
 - Order data, order payment, order review, order items
 - Product, product translation

OVERVIEW & TARGETS

Datasets overview

	GEOLOCATION				
	geolocation_lat				
	geolocation_lng				
	geolocation_zip_code_prefix				
	geolocation_city				
	SELLER				
	seller_id				
	seller_zip_code_prefix				
	seller_city				
	seller_state				
	CUSTOMER				
	customer_id				
	customer_unique_id				
	customer_zip_code_prefix				
	customer_city				
	ORDER DATA				
	order_id				
	customer_id				
	order_status				
	order_purchase_timestamp				
	ORDER ITEMS				
	order_id				
	order_item_id				
	product_id				
	seller_id				
	PRODUCT				
	product_id				
	product_category_name				
	product_name_lenght				
	product_description_lenght				
	PRODUCT TRANSLATION				
	product_category_name				
	product_category_name_english				
	ORDER REVIEW				
	review_id				
	order_id				
	review_score				
	review_comment_title				
	ORDER PAYMENT				
	order_id				
	payment_sequential				
	payment_type				
	payment_installments				
	PRODUCT				
	product_weight_g				
	product_length_cm				
	product_height_cm				
	product_width_cm				
	ORDER REVIEW				
	review_creation_date				
	review_answer_timestamp				
	ORDER PAYMENT				
	payment_value				

OVERVIEW & TARGETS



02



DATA MERGING & ANALYSIS



DATA MERGING

Between products, and their translation

- Missing categories added under a translation dictionary
- Inner merge between products, and category translation, on the 'product_category_name' feature.
- Dropped portuguese version
- Regrouped categories into 9 large sets
 - hygiene, electronics, furniture, leisure, fashion, groceries, office, diy, misc
- Added missing dimensions and weight by category average
- Calculated volume from dimensions
- Removed:
 - name length, description, photos, length, height, width

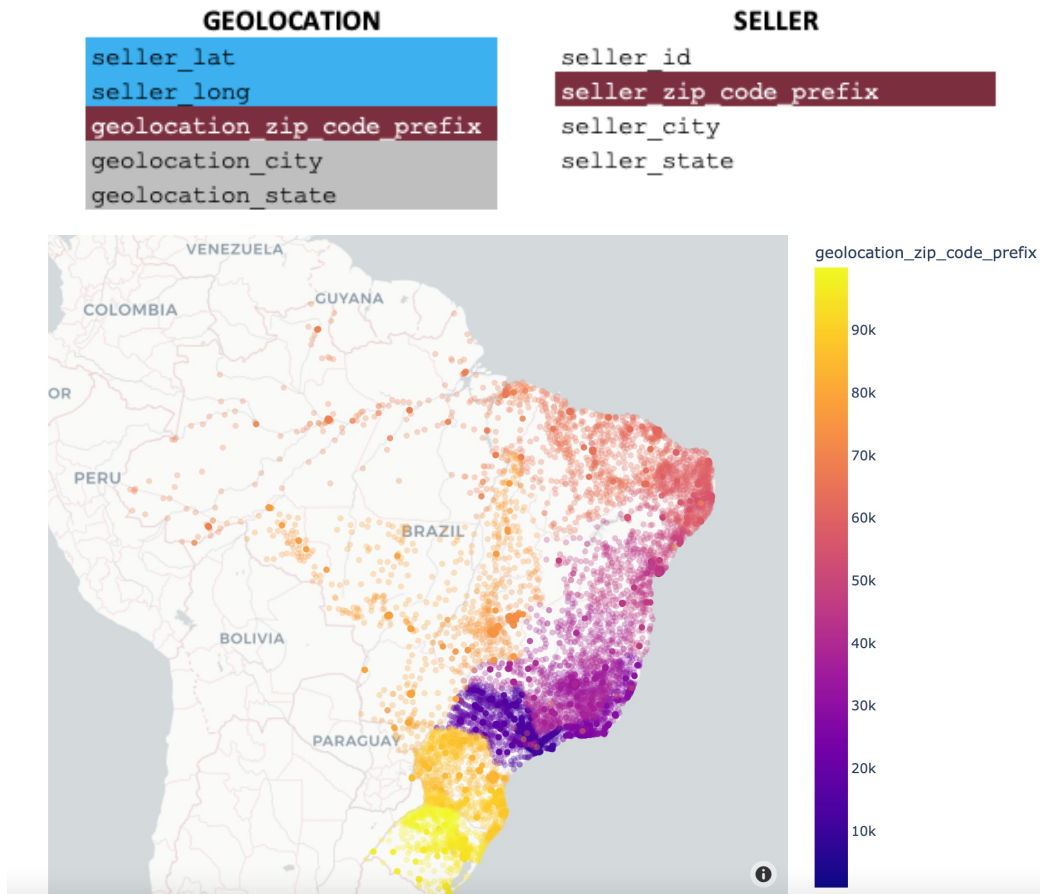
PRODUCT	
product_id	
product_category_name	
product_name_lenght	
product_description_lenght	
product_photos_qty	
product_weight_g	
product_length_cm	
product_height_cm	
product_width_cm	
product_volume_cm3	

PRODUCT TRANSLATION	
product_category_name	
product_category_name_english	

DATA MERGING

Between sellers, and their location

- Dropped duplicates of ZIP code
- Removed geographical outliers based on Brazil coordinates
 - $-35 < \text{latitude} < 5$
 - $-75 < \text{longitude} < -35$
- Added sellers location by Left join on seller dataset, on zip code
- State average coordinate to fill missing coordinates
- Drop: 'state', 'city', 'zip codes'



DATA MERGING

Order items with products and sellers

- Left join on order items from products, based on 'product_id' feature
- As there are multiple items per order, these items are aggregated for each order.
- Left join on order items from sellers, on 'seller_id'

```
agg_dict = {'order_item_id': np.max,  
            'seller_id': np.max,  
            'price': np.sum,  
            'freight_value': np.sum,  
            'product_weight_g': np.mean,  
            'product_volume': np.mean,  
            'product_category_name_english': mode  
            }
```

SELLER

```
seller_id  
seller_zip_code_prefix  
seller_city  
seller_state  
seller_lat  
seller_long
```

ORDER ITEMS

```
order_id  
order_item_id  
product_id  
seller_id  
shipping_limit_date  
price  
freight_value
```

PRODUCT

```
product_id  
product_category_name_english  
product_weight_g  
product_volume_cm3
```

DATA MERGING

Regrouping all orders data

- Order payments rows with same order id are grouped
- Only orders that are already delivered are considered, the others are discarded
- Left join of order_dataset with order_item, on 'order_id'
- Left join of orders with order_reiew, on 'order_id'
- Left join of orders with oder_payments, on 'order_id'
- Missing review scores are replaced by the median grade

3

```
agg_dict2 = {  
    'perc_credit': np.mean,  
    'payment_installments': np.mean,  
    'payment_value': np.sum  
}
```

ORDER DATA

order_id
customer_id
order_status
order_purchase_timestamp
order_approved_at
order_delivered_carrier_date
order_delivered_customer_date
order_estimated_delivery_date

ORDER ITEMS

order_id
order_item_id
product_id
seller_id
shipping_limit_date
price
freight_value
seller_city
seller_state
seller_lat
seller_long
product_category_name_english
product_weight_g
product_volume_cm3

ORDER REVIEW

review_id
order_id
review_score
review_comment_title
review_comment_message
review_creation_date
review_answer_timestamp

ORDER PAYMENT

order_id
payment_sequential
perc_credit
payment_installments
payment_value

DATA MERGING

Regrouping all customer data

- Left join of customers with orders, on 'customer_id'
- 2972 customers have no order, these rows are dropped
- Left join of customer with geolocation, on zip code
- Filling missing customer coordinates by using states

CUSTOMER

```
customer_id  
customer_unique_id  
customer_zip_code_prefix  
customer_city  
customer_state
```

GEOLOCATION

```
geolocation_lat  
geolocation_lng  
geolocation_zip_code_prefix  
geolocation_city  
geolocation_state
```

ORDER DATA

```
customer_id  
order_purchase_timestamp  
order_delivered_customer_date  
order_item_id  
price  
freight_value  
seller_city  
seller_state  
seller_lat  
seller_long  
product_category_name_english  
product_weight_g  
product_volume_cm3  
review_score  
perc_credit  
payment_installments  
payment_value
```

DATA CLEANING

Duplicates

- Customer_id, and order_id show 529 duplicates. They are removed.
- 'customer_id' feature is removed, as it is unique to order_id, and its name is misleading

Location

- Customer_city contains hundreds of cities, which don't include significant part of the population. This feature is removed.
- States are grouped in 5 regions. 2 regions have more than 95% orders.

```
state_dict = {  
    'North': ['AC', 'AP', 'AM', 'PA', 'RO', 'RR', 'TO'],  
    'Northeast': ['AL', 'BA', 'CE', 'MA', 'PB', 'PE', 'PI',  
                  'RN', 'SE'],  
    'Southeast': ['ES', 'MG', 'RJ', 'SP'],  
    'South': ['PR', 'RS', 'SC'],  
    'Centerwest': ['DF', 'GO', 'MT', 'MS']  
}
```

DATA CLEANING

Payment

- 'Price' feature can be removed as payment value already gives the total.

Dates

- Purchase and delivery dates are converted with `pd.to_datetime`
- 'delivery_time' is created, it's the difference between date of delivery and date of purchase.

DATA CLEANING

Customers aggregation

- Aggregation is carried out while grouping by customer_unique_id

```
agg_dict = {  
    'nb_of_orders': np.max,  
    'customer_state': mode,  
    'freight_value': np.sum,  
    'perc_credit': np.mean,  
    'payment_installments': np.mean,  
    'payment_value': np.sum,  
    'category': mode,  
    'product_weight_g': np.mean,  
    'product_volume': np.mean,  
    'nb_of_items': np.mean,  
    'review_score': np.mean,  
    'date_purchase': np.max,  
    'delivery_time': np.mean,  
}
```

- Recency is calculated from date of purchase
- Frequency is directly given by nb of orders
- Monetary is directly taken from payment_value

DATA CLEANING

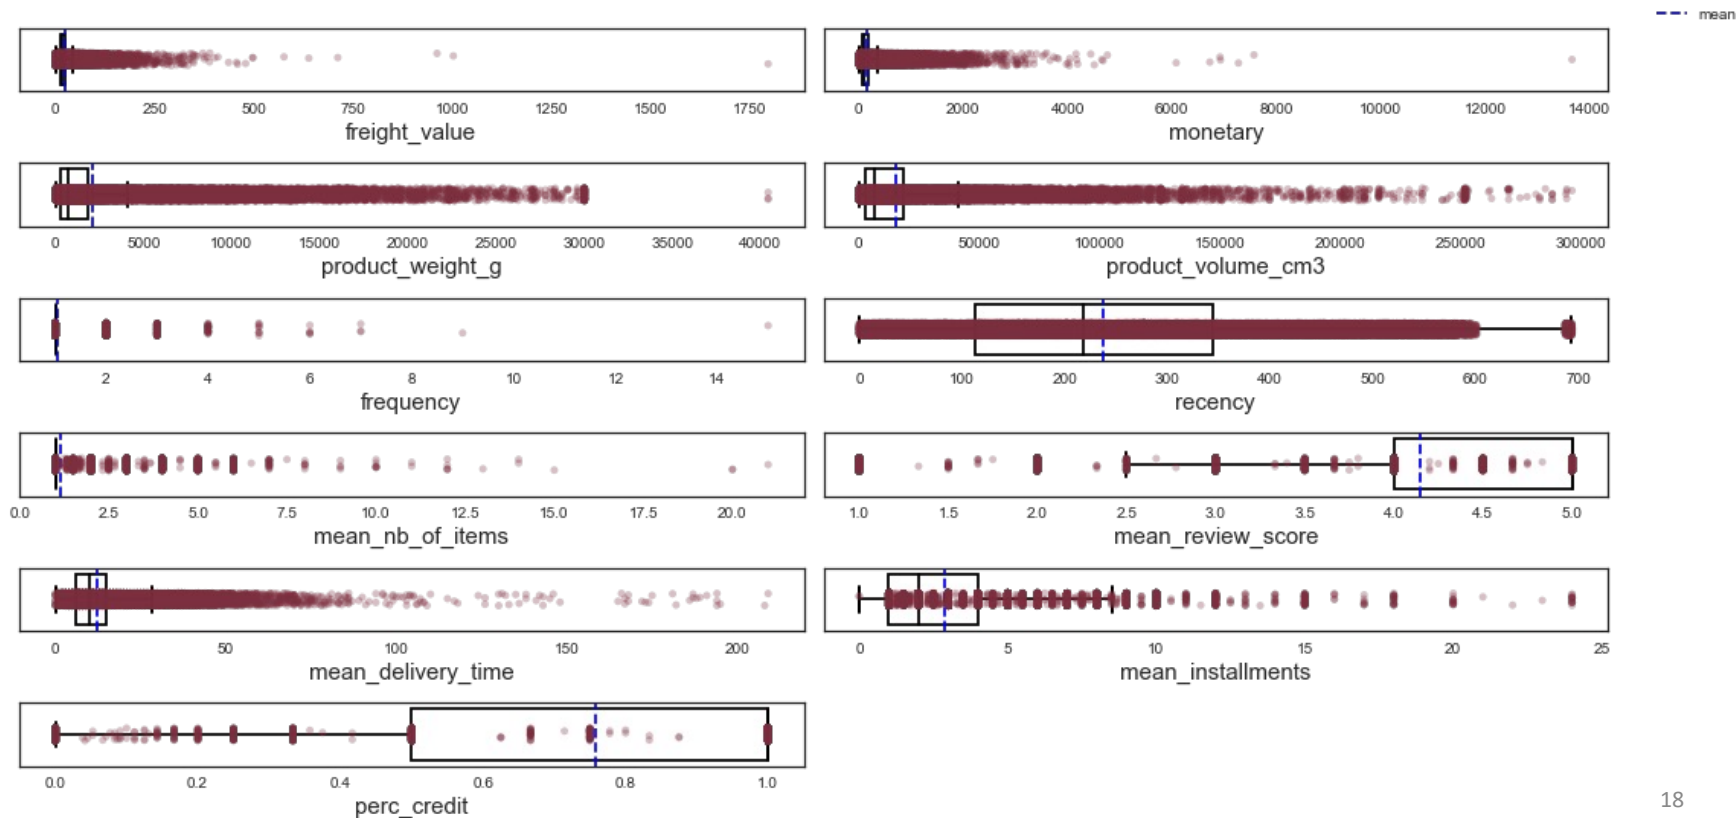
Cleaned features overview

Customer information	ID customer_state
Order information	frequency product_weight_g product_volume_cm3 category mean_nb_of_items
Time	mean_delivery_time date_purchase recency
Payment	perc_credit mean_installments monetary freight_value
Review	mean_review_score

EXPLORATORY ANALYSIS

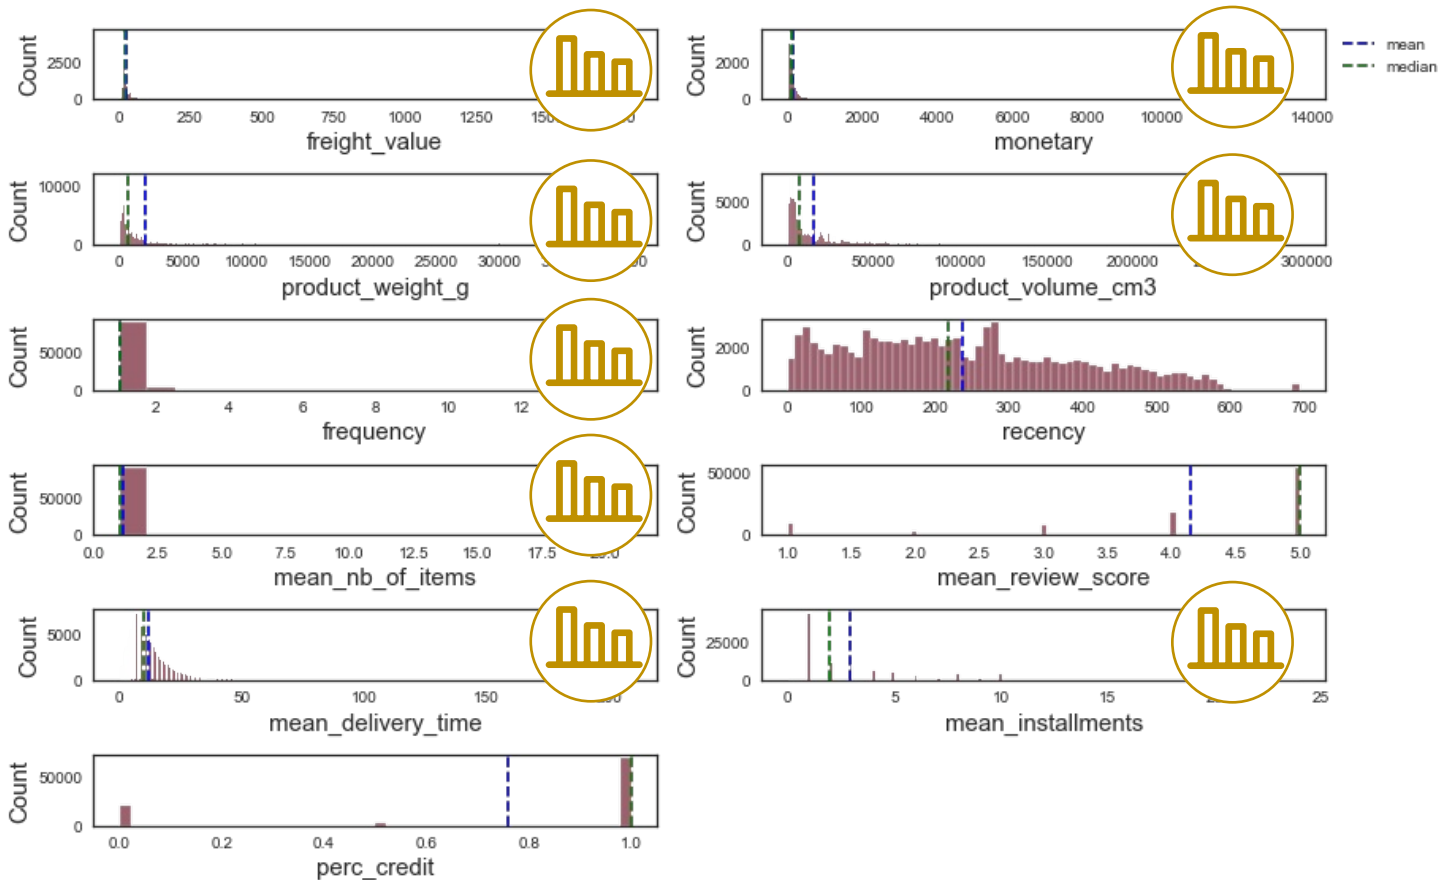
Numerical features

Quantitative variables distribution



EXPLORATORY ANALYSIS

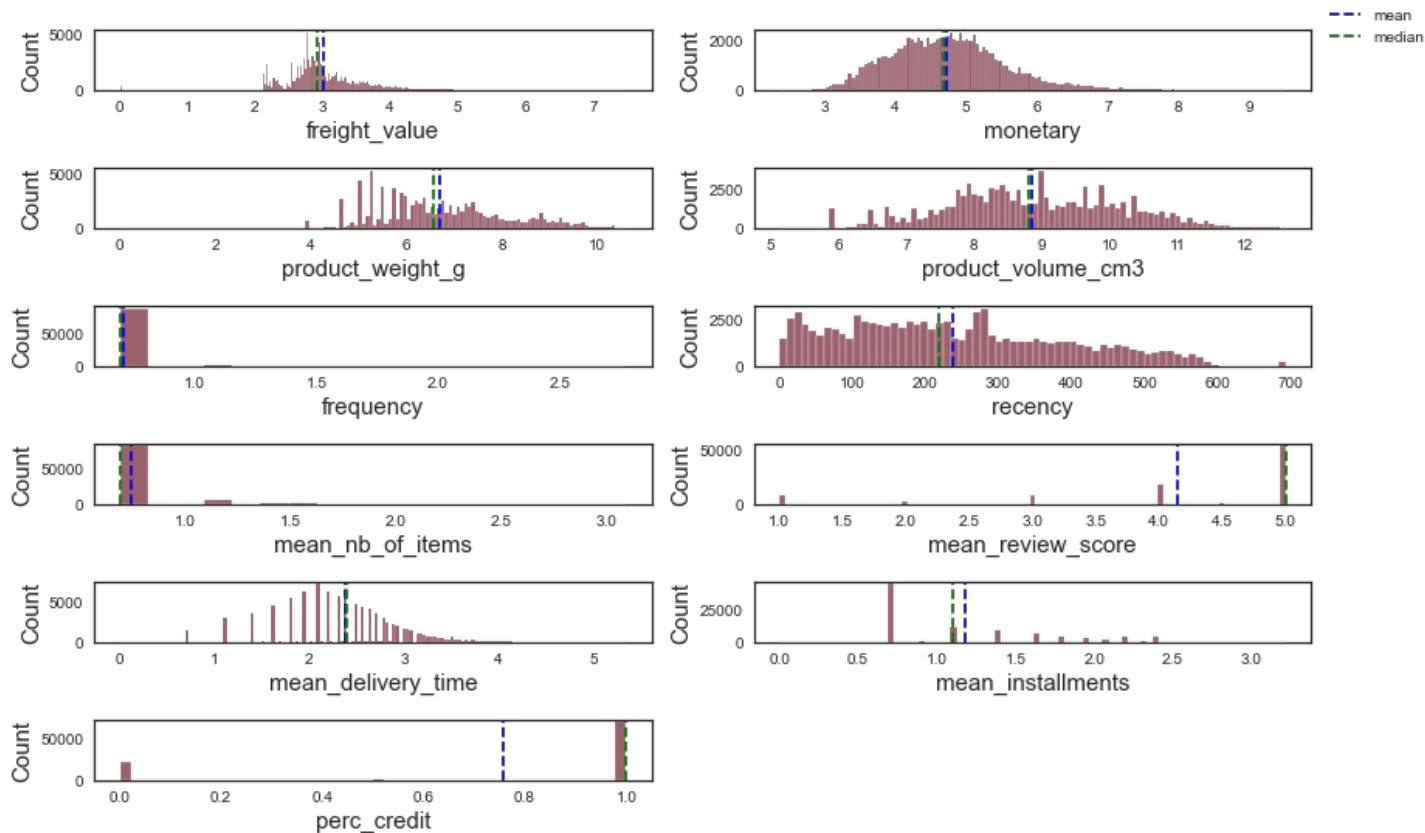
Numerical variables



EXPLORATORY ANALYSIS

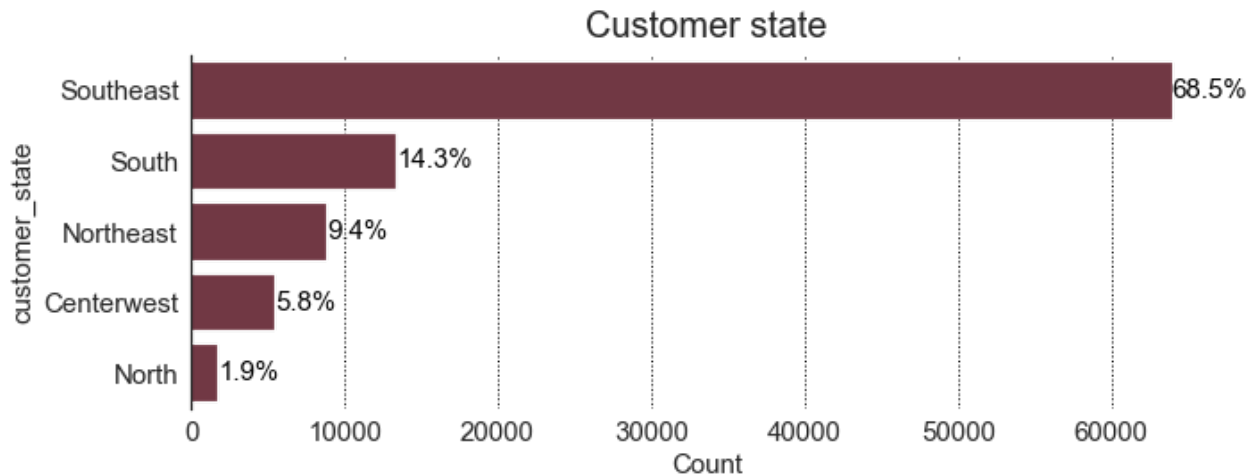
Numerical variables

Log distribution of quantitative features




EXPLORATORY ANALYSIS

Categorical features



EXPLORATORY ANALYSIS

Categorical features




EXPLORATORY ANALYSIS



EXPLORATORY ANALYSIS

Multivariate analysis

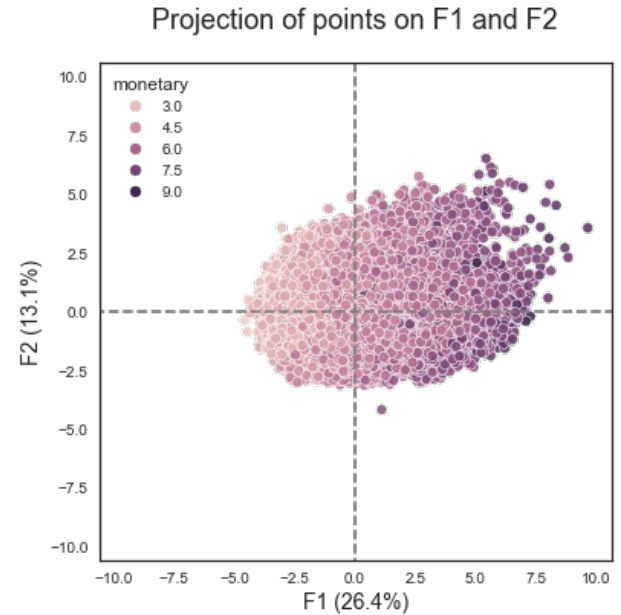
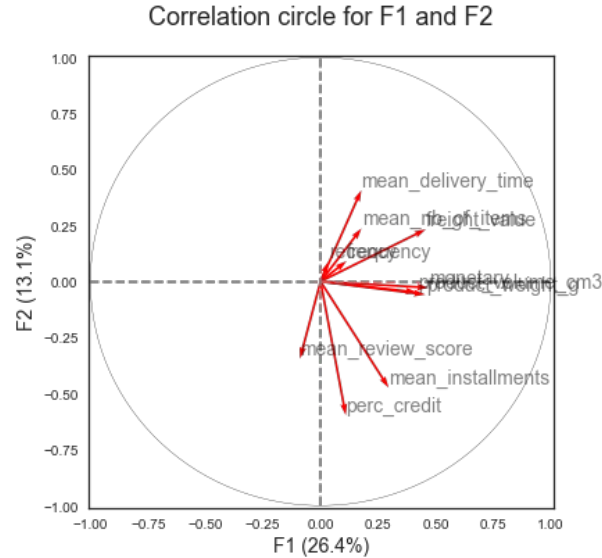

Correlation heatmap



EXPLORATORY ANALYSIS

Multivariate analysis

Principal Component Analysis



EXPLORATORY ANALYSIS

Multivariate analysis


ANOVA

$$\eta^2 = \frac{ESS}{TSS}$$

	Numerical_features	customer_state	category
0	freight_value	0.112305	0.026310
1	monetary	0.013882	0.022516
2	product_weight_g	0.000097	0.183267
3	product_volume_cm3	0.000592	0.235835
4	frequency	0.000124	0.004269
5	recency	0.001761	0.011464
6	mean_nb_of_items	0.000380	0.011825
7	mean_review_score	0.002563	0.002841
8	mean_delivery_time	0.142176	0.006026
9	mean_installments	0.005778	0.023064
10	perc_credit	0.001153	0.002137

03

CLUSTERING



CLUSTERING

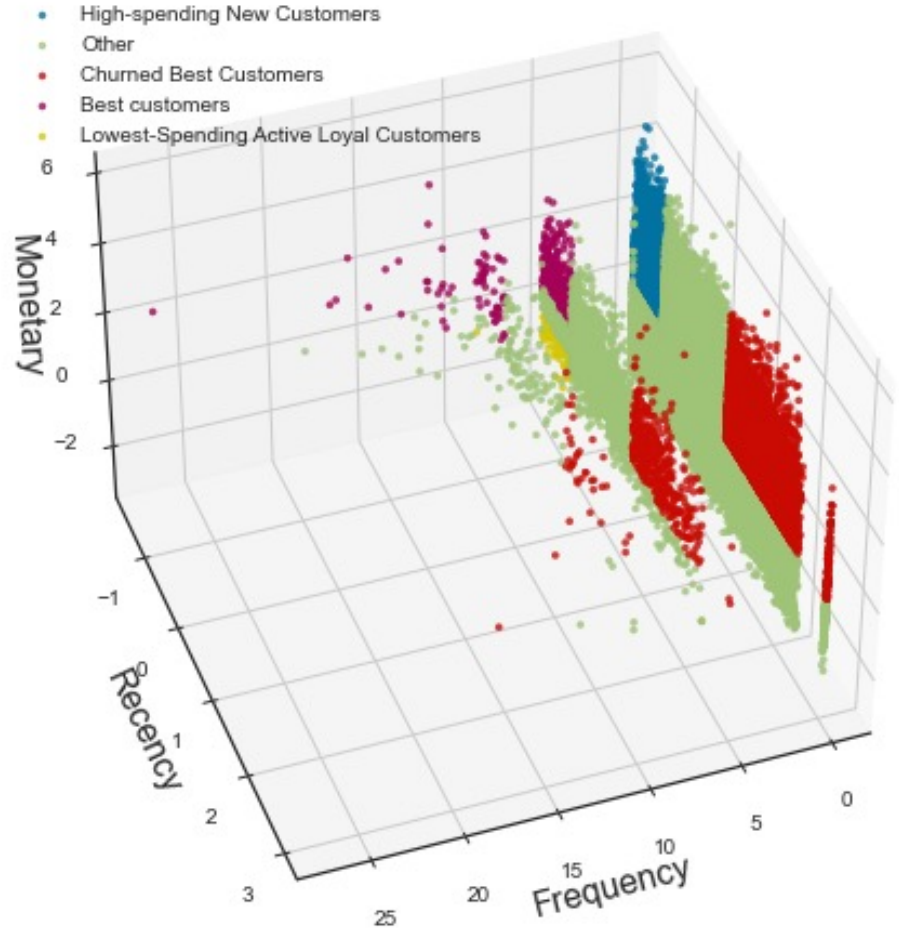

RFM Segmentation

- Recency, Frequency, Monetary features are selected
- They are normalized through StandardScaler
- Segmentation is carried out based on tiers:
 - 4 tiers for recency and monetary
 - Only 2 tiers for frequency due to limited nb of customers that come more than once
- Names are attributed based on scores:
 - 1-1-1 Best customers
 - 1-2-1 and 1-2-2: High-spending new customers
 - 1-1-3 and 1-1-4: Lowest-spending active loyal customers
 - 4-1-1, 4-1-2, 4-2-1 and 4-2-2: Churned best customers

CLUSTERING

RFM Segmentation

- Low number of customers coming more than once limit the relevance of this type of study.



CLUSTERING

RFM Segmentation

- Recency, Frequency, Monetary features are selected
- They are normalized through StandardScaler
- Segmentation is carried out based on tiers:
 - 4 tiers for recency and monetary
 - Only 2 tiers for frequency due to limited nb of customers that come more than once
- Names are attributed based on scores:
 - 1-1-1 Best customers
 - 1-2-1 and 1-2-2: High-spending new customers
 - 1-1-3 and 1-1-4: Lowest-spending active loyal customers
 - 4-1-1, 4-1-2, 4-2-1 and 4-2-2: Churned best customers

CLUSTERING

Selection of more features as an alternative to RFM


- A random sample of 10000 is considered, to reduce convergency time.
- Follow features are selected:
 - Monetary
 - Frequency
 - Recency
 - Percentage of credit
 - Mean review score
 - Mean number of items
 - Mean product weight
- Features are normalized with StandardScaler

CLUSTERING

Selection of more features as an alternative to RFM

K-means clustering

- Elbow method based on distortion gives a elbow at $k = 7$ clusters




CLUSTERING

Selection of more features as an alternative to RFM

K-means clustering

- Metrics also confirm 7 clusters is an interesting choice




CLUSTERING

K-means clustering

Selection of more features as an alternative to RFM

Clustering visualization




CLUSTERING

Selection of more features as an alternative to RFM

Hierarchical clustering

- Elbow method based on distortion gives a elbow at $k = 6$ clusters




CLUSTERING

Selection of more features as an alternative to RFM

Hierarchical clustering

- Metrics also confirm 6 clusters is an interesting choice, mostly for stability




CLUSTERING

Selection of more features as an alternative to RFM

Hierarchical clustering


Clustering visualization



CLUSTERING

Selection of more features as an alternative to RFM


Hierarchical clustering



CLUSTERING

Selection of more features as an alternative to RFM

DBScan clustering




CLUSTERING

Selection of more features as an alternative to RFM

DBScan clustering


- Metrics confirm $\text{eps} = 4$ gives best results for silhouette



CLUSTERING

Selection of more features as an alternative to RFM

DBScan clustering



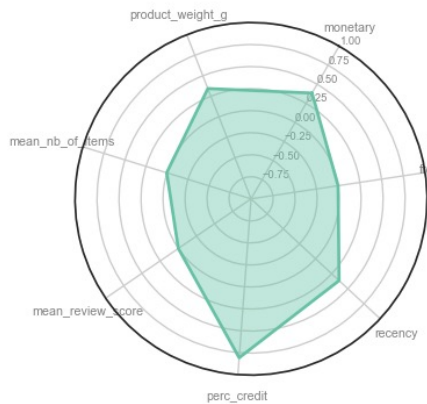
CLUSTERING

Selection of more features as an alternative to RFM


Best model and radar plot

- K-means has the best stability according to ARI score, it is kept with 7 clusters


Unsatisfied




New credit paying satisfied



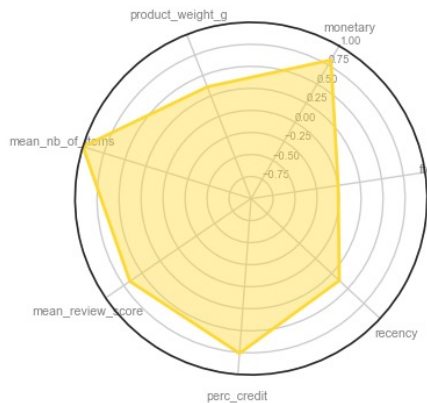
Cash paying casual satisfied




Compulsive, few expensive items




Compulsive, many items



High value frequent buyer, to be reached




Credit paying casual satisfied



MAINTENANCE

K-means clustering with 7 clusters

- ARI score as a function of time
- There seems to be a drop after 45 days



04



CONCLUSION

CONCLUSION

- Data available from October 2016 to October 2018
- 9 Dataframes for a total of 99k customers id
 - Customer, Seller, Geolocation
 - Order data, order payment, order review, order items
 - Product, product translation
- Objective: cluster customers with unsupervised learning
- Only 3% of customers buy more than once, RFM is therefore limited
- Best clustering model: K-means, with 7 clusters
- Maintenance proposed: every 45 days, based on stability



THANK YOU

Victor Benard